

Image Processing - *American Sign Language (ASL)*



Group Members :

Name	Roll No
R.Abhinav	CB.EN.U4CSE19453
P.Kalaiarasan	CB.EN.U4CSE19446
P.Koushik	CB.EN.U4CSE19449
S.Shanthan	CB.EN.U4CSE19459

Review-1

Mounting and Importing

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).

All the import statements for the libraries used in the project below lies here.

```
In [ ]: #Importing Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import random
import warnings
warnings.filterwarnings('ignore')
```

Exploring the Data

```
In [ ]: #reading the dataset
```

```
test_df = pd.read_csv("/content/drive/MyDrive/Sign Language/datasets/sign_mnist_test.csv")
train_df = pd.read_csv("/content/drive/MyDrive/Sign Language/datasets/sign_mnist_train.csv")
```

```
In [ ]: train_df.tail()
```

Out[]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	pixel10	pixel11
27450	13	189	189	190	190	192	193	193	193	193	194	195
27451	23	151	154	157	158	160	161	163	164	166	167	168
27452	18	174	174	174	174	174	175	175	174	173	173	174
27453	17	177	181	184	185	187	189	190	191	191	190	191
27454	23	179	180	180	180	182	181	182	183	182	182	183

5 rows × 785 columns

◀ ▶

```
In [ ]: print(type(test_df))
```

```
<class 'pandas.core.frame.DataFrame'>
```

In []: `test_df.head()`

Out[]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	pixel10	pixel11	pi
0	6	149	149	150	150	150	151	151	150	151	152	152	
1	5	126	128	131	132	133	134	135	135	136	138	137	
2	10	85	88	92	96	105	123	135	143	147	152	157	
3	0	203	205	207	206	207	209	210	209	210	209	208	
4	3	188	191	193	195	199	201	202	203	203	203	204	

5 rows × 785 columns

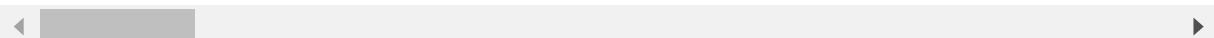


In []: `test_df.tail()`

Out[]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	pixel10	pixel11
7167	1	135	119	108	102	105	99	61	103	121	133	143
7168	12	157	159	161	164	166	166	171	174	175	176	176
7169	2	190	191	190	191	190	190	192	192	191	192	193
7170	4	201	205	208	209	214	216	218	223	226	229	234
7171	2	173	174	173	174	173	173	175	175	174	175	176

5 rows × 785 columns



In []: `print(test_df.info())`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7172 entries, 0 to 7171
Columns: 785 entries, label to pixel784
dtypes: int64(785)
memory usage: 43.0 MB
None
```

In []: # Showing the first 4 lines of the training set

`train_df.iloc[:4,:10]`

Out[]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9
0	3	107	118	127	134	139	143	146	150	153
1	6	155	157	156	156	156	157	156	158	158
2	2	187	188	188	187	187	186	187	188	187
3	2	211	211	212	212	211	210	211	210	210

```
In [ ]: print(f'Number of images in the training set: {train_df.shape[0]}')
print(f'Number of images in the test set: {test_df.shape[0]}')

d = int((train_df.shape[1] - 1)**0.5)
print(f'Shape of the images: {d} x {d}' )
```

Number of images in the training set: 27455
 Number of images in the test set: 7172
 Shape of the images: 28 x 28

```
In [ ]: print(train_df.dtypes)
print("_____ \n")
print(train_df.dtypes.value_counts())
```

label	int64
pixel1	int64
pixel2	int64
pixel3	int64
pixel4	int64
	...
pixel780	int64
pixel781	int64
pixel782	int64
pixel783	int64
pixel784	int64
Length:	785, dtype: object

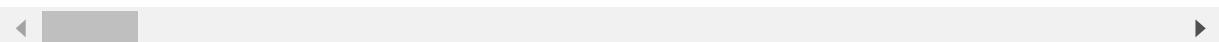
int64	785
dtype:	int64

```
In [ ]: train_df.describe()
```

Out[]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	
count	27455.000000	27455.000000	27455.000000	27455.000000	27455.000000	27455.000000	274
mean	12.318813	145.419377	148.500273	151.247714	153.546531	156.210891	1
std	7.287552	41.358555	39.942152	39.056286	38.595247	37.111165	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	6.000000	121.000000	126.000000	130.000000	133.000000	137.000000	1
50%	13.000000	150.000000	153.000000	156.000000	158.000000	160.000000	1
75%	19.000000	174.000000	176.000000	178.000000	179.000000	181.000000	1
max	24.000000	255.000000	255.000000	255.000000	255.000000	255.000000	2

8 rows × 785 columns



Statistical Analysis of data

```
In [ ]: train_df.mean()
```

```
Out[ ]: label      12.318813
         pixel1     145.419377
         pixel2     148.500273
         pixel3     151.247714
         pixel4     153.546531
         ...
         pixel780    162.736696
         pixel781    162.906137
         pixel782    161.966454
         pixel783    161.137898
         pixel784    159.824731
Length: 785, dtype: float64
```

```
In [ ]: #finding the median for the numerical columns in the data set
train_df.median()
```

```
Out[ ]: label      13.0
         pixel1     150.0
         pixel2     153.0
         pixel3     156.0
         pixel4     158.0
         ...
         pixel780    184.0
         pixel781    184.0
         pixel782    182.0
         pixel783    182.0
         pixel784    182.0
Length: 785, dtype: float64
```

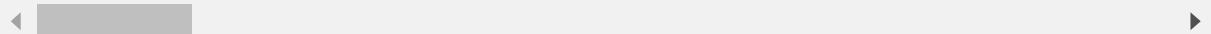
```
In [ ]: # finding the mode of the numerical columns in the data set
```

```
k = train_df.mode()
k[0:1]
```

```
Out[ ]:
```

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	pixel10	pixel11	pi
0	17.0	158.0	171.0	172.0	171.0	165.0	159.0	157.0	158.0	174.0	164.0	165.0	165.0

1 rows × 785 columns



```
In [ ]: # finding the standard deviation the dataset for the numerical values  
standard_deviations = train_df.std()  
print(standard_deviations)
```

```
label      7.287552  
pixel1     41.358555  
pixel2     39.942152  
pixel3     39.056286  
pixel4     38.595247  
...  
pixel780    63.444008  
pixel781    63.509210  
pixel782    63.298721  
pixel783    63.610415  
pixel784    64.396846  
Length: 785, dtype: float64
```

```
In [ ]: # finding the variance for the numerical values of the dataset  
variance = standard_deviations**2  
print(variance)
```

```
label      53.108415  
pixel1    1710.530097  
pixel2    1595.375528  
pixel3    1525.393469  
pixel4    1489.593076  
...  
pixel780    4025.142186  
pixel781    4033.419724  
pixel782    4006.728058  
pixel783    4046.284910  
pixel784    4146.953763  
Length: 785, dtype: float64
```

Data preprocessing

- converting array to images(tensors)
- visualizing lables and making sure that dataset is balanced
- performing encoding for lables

In []: # couting the missing values or cells with empty values

```
no_of_missing_values = train_df.isnull().sum()
total_missing_values = no_of_missing_values.sum()
print(no_of_missing_values)

print("\n\nTotal missing values in the dataset: {0}".format(total_missing_values))
```

```
label      0
pixel1     0
pixel2     0
pixel3     0
pixel4     0
...
pixel780   0
pixel781   0
pixel782   0
pixel783   0
pixel784   0
Length: 785, dtype: int64
```

Total missing values in the dataset: 0

In []: # checking if there are any duplicated values

```
print("Duplicated Columns : ",train_df.duplicated().sum())
```

Duplicated Columns : 0

In []: train_df

Out[]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	pixel10	pixel11
0	3	107	118	127	134	139	143	146	150	153	156	158
1	6	155	157	156	156	156	157	156	158	158	157	158
2	2	187	188	188	187	187	186	187	188	187	186	188
3	2	211	211	212	212	211	210	211	210	210	211	209
4	13	164	167	170	172	176	179	180	184	185	186	188
...
27450	13	189	189	190	190	192	193	193	193	193	194	193
27451	23	151	154	157	158	160	161	163	164	166	167	168
27452	18	174	174	174	174	174	175	175	174	173	173	174
27453	17	177	181	184	185	187	189	190	191	191	190	191
27454	23	179	180	180	180	182	181	182	183	182	182	182

27455 rows × 785 columns



```
In [ ]: # Create training and testing arrays
train_set = np.array(train_df, dtype = 'float32')
test_set = np.array(test_df, dtype='float32')
```

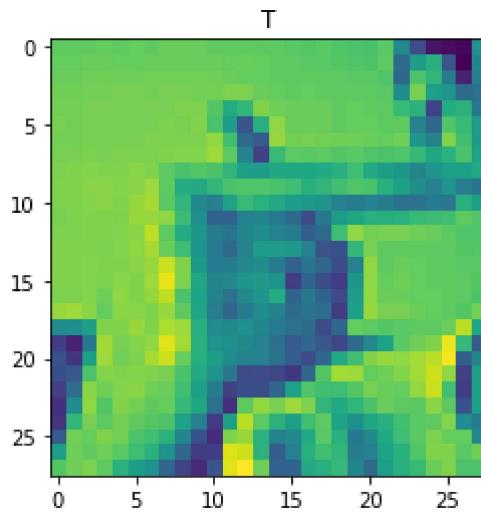
Reason:

- The training and testing input arrays are converted to continuous float values since it allows our model for a more precise learning as compared to discrete values

```
In [ ]: #Specifying class Labels
class_names = ['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M',
, 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y' ]
```

```
In [ ]: #See a random image for class label verification
import random
i = random.randint(1,27455)
plt.imshow(train_set[i,1:].reshape((28,28)))
label_index = train_df["label"][i]
plt.title(f"{class_names[label_index]}")
```

Out[]: Text(0.5, 1.0, 'T')

**Inference:**

- We also need to reshape the array to (28x28) since the initial shape is just a row array.

Subplots :

- Define the dimensions of the plot grid
- flatten the 15×15 matrix into 225 array
- get the length of the train dataset
- Select a random number from 0 to `n_train`
- create evenly spaces variables
- Select a random number
- read and display an image with the selected index

```
In [ ]: #
W_grid = 5
L_grid = 5
fig, axes = plt.subplots(L_grid, W_grid, figsize = (10,10))
axes = axes.ravel()
n_train = len(train_set)
for i in np.arange(0, W_grid * L_grid):
    index = np.random.randint(0, n_train)
    axes[i].imshow( train_set[index,1:].reshape((28,28)) )
    label_index = int(train_set[index,0])
    axes[i].set_title(class_names[label_index], fontsize = 8)
    axes[i].axis('off')
plt.subplots_adjust(hspace=0.4)
```



```
In [ ]: train_x = train_df[train_df.columns[1::]].to_numpy()
train_y = train_df[train_df.columns[0]].to_numpy()

test_x = test_df[test_df.columns[1::]].to_numpy()
test_y = test_df[test_df.columns[0]].to_numpy()

print("SUMMARY OF DATA:")

print("train_x shape: " + str(train_x.shape))
print("train_y shape: " + str(train_y.shape))
print("test_x shape: " + str(test_x.shape))
print("test_y shape: " + str(test_y.shape))
```

```
SUMMARY OF DATA:
train_x shape: (27455, 784)
train_y shape: (27455,)
test_x shape: (7172, 784)
test_y shape: (7172,)
```

- we know that the pixel values lies between 0-255 but it is observed that models performs exceptionally well if we scale pixel values between 0-1

Visualization

```
In [ ]: #Visualize train images
plt.figure(figsize=(10, 10))
for i in range(25):
    plt.subplot(5, 5, i + 1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(train_x[i].reshape((28,28)), cmap=plt.cm.binary)
    label_index = int(train_y[i])
    plt.title(class_names[label_index])
plt.show()
```



```
In [ ]: #Get our training Labels:
labels = train_df['label'].values
```

```
In [ ]: label_frq=pd.value_counts(train_df['label'],ascending=True).reset_index(level=0)
pd.DataFrame(label_frq)
```

Out[]:

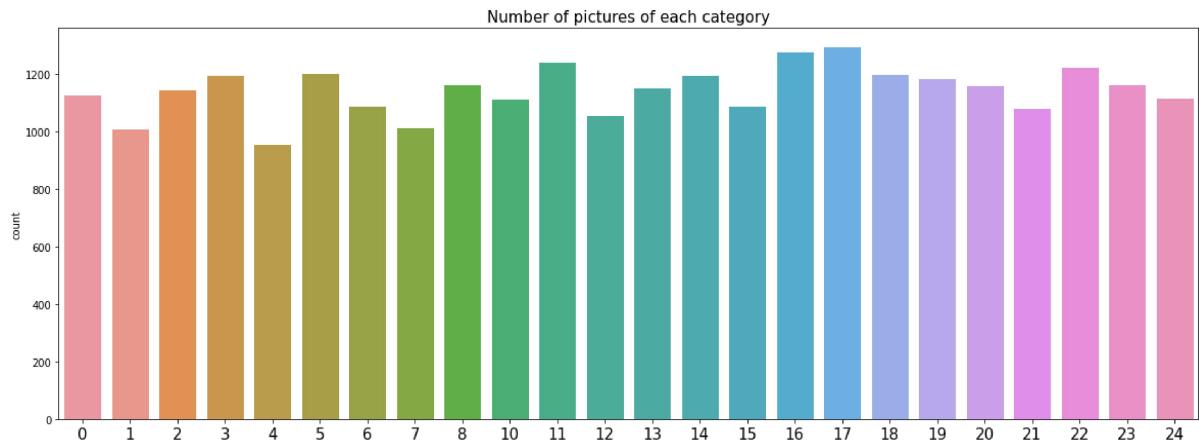
	index	label
0	4	957
1	1	1010
2	7	1013
3	12	1055
4	21	1082
5	15	1088
6	6	1090
7	10	1114
8	24	1118
9	0	1126
10	2	1144
11	13	1151
12	20	1161
13	8	1162
14	23	1164
15	19	1186
16	3	1196
17	14	1196
18	18	1199
19	5	1204
20	22	1225
21	11	1241
22	16	1279
23	17	1294

```
In [ ]: unique_labels = train_df['label'].unique()
unique_labels = np.sort(unique_labels)
unique_labels
```

```
Out[ ]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8, 10, 11, 12, 13, 14, 15, 16, 17,
   18, 19, 20, 21, 22, 23, 24])
```

In []: # visualizing the data

```
plt.figure(figsize=(20,7))
plt.title("Number of pictures of each category", fontsize = 15)
plt.xticks(fontsize = 15)
sns.set_style("darkgrid");
sns.countplot(train_y);
```



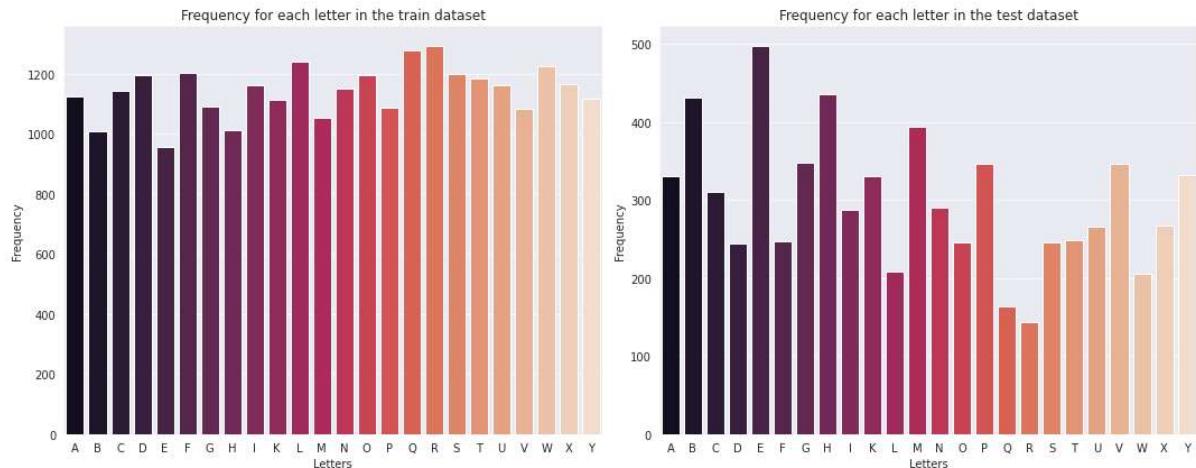
Observation:

- We can see the visual tells us that there are nearly 1000 examples for each class output
- So we can consider this dataset as a balanced dataset because there's no class suffering from very less or too much examples

```
In [ ]: list_data = [train_df, test_df]
fig,axes = plt.subplots(nrows=1, ncols=2, figsize=(15,6))

for data, ax, names in zip(list_data, axes.ravel(), ['train', 'test']):
    sns.countplot(data['label'], palette='rocket', ax=ax)
    ax.set_title("Frequency for each letter in the {} dataset".format(names))
    ax.set_xlabel('Letters')
    ax.set_ylabel('Frequency')
    ax.set_xticklabels(['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'K', 'L', 'M', 'N',
'P', 'Q', 'R', 'S',
'T', 'U', 'V', 'W', 'X', 'Y'])

plt.tight_layout()
```



```
In [ ]: #Encoding:
```

```
from sklearn.preprocessing import LabelBinarizer
labels = train_df['label'].values
label_binarizer = LabelBinarizer()
labels = label_binarizer.fit_transform(labels)
labels
```

```
Out[ ]: array([[0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 1, ..., 0, 0, 0],
 ...,
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 1, 0]])
```

Observation:

- So basically LabelBinarizer performed One-Hot Encoding on our training data.

```
In [ ]: labels[100]
```

```
Out[ ]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
In [ ]: import cv2
df_htrain = pd.read_csv("/content/drive/MyDrive/Sign Language/Hog_Features.csv")
df_htest = pd.read_csv("/content/drive/MyDrive/Sign Language/Hog_Features_test.csv")
htrain = df_htrain[df_htrain.columns[1::]].to_numpy()
htest = df_htest[df_htest.columns[1::]].to_numpy()
print("Training Data size : ",htrain.shape)
print("Test Data size : ",htest.shape)
```

```
Training Data size : (27455, 1764)
Test Data size : (7172, 1764)
```

Review-2(Final Review)

Feature Engineering

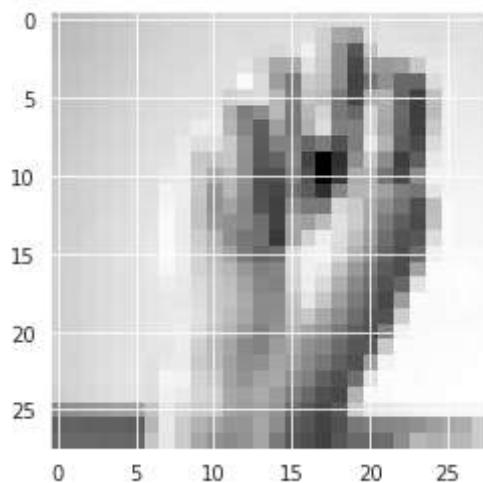
```
In [ ]: train_df_original = train_df.copy()
```

```
In [ ]: train_x = train_x/255
test_x = test_x/255
```

```
In [ ]: # Split into training, test and validation sets
val_index = int(train_df.shape[0]*0.1)

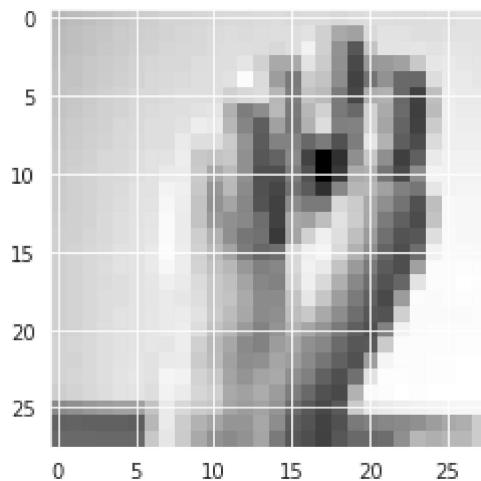
train_df = train_df_original.iloc[val_index:]
val_df = train_df_original.iloc[:val_index]
```

```
In [ ]: index = 4
plt.imshow(train_x[index].reshape(28, 28),cmap='gray')
plt.show()
```



```
In [ ]: import matplotlib.pyplot as plt
from skimage.feature import hog
from skimage import data,exposure
import cv2
from sklearn import svm
%matplotlib inline
np.random.seed(1)
```

```
In [ ]: image = plt.imshow(train_x[index].reshape(28, 28),cmap='gray')
```



End of the Review😊

In this work, we propose a method to detect American Sign Language by analyzing sign language MNIST image dataset. We will use the pixel level features in each image as well as HOG Features, to be followed with feature selection to improve the performance. Finally we will print different evaluation metrics for the classifier model such as Accuracy, F1 Score, Precision, Recall and plot ROC curves for all the classes.

Feature Extraction Using HOG Features

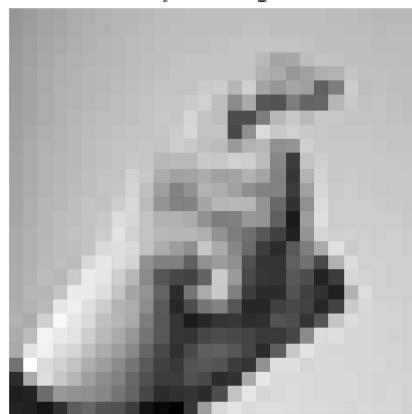
```
In [ ]: from skimage.feature import hog
from skimage import data, exposure
def get_hog(image):
    fd,hog_image=hog(image.reshape(28,28),pixels_per_cell=(2,2),
                      cells_per_block=(1, 1),visualize=True,feature_vector=F
else)
    return(fd,hog_image)

def show_hog(image,hog_image):
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(8, 4), sharex=True, sharey=T
rue)
    ax1.axis('off')
    ax1.imshow(image.reshape(28, 28), cmap=plt.cm.gray)
    ax1.set_title('Input image')
    hog_image_rescaled = exposure.rescale_intensity(hog_image, in_range=(0, 10
))

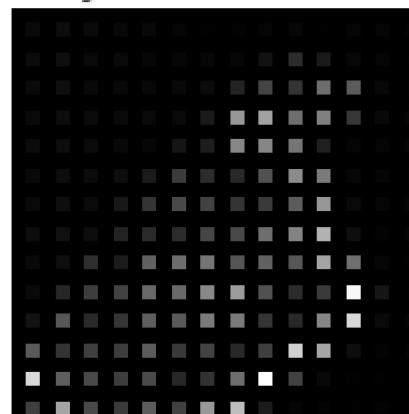
    ax2.axis('off')
    ax2.imshow(hog_image_rescaled, cmap=plt.cm.gray)
    ax2.set_title('Histogram of Oriented Gradients')
    plt.show()
```

```
In [ ]: a,b=get_hog(train_x[180])
show_hog(train_x[180],b)
a=a.reshape(-1)
a=a.reshape(1,len(a))
print(a.shape)
```

Input image



Histogram of Oriented Gradients



(1, 1764)

Feature Selection using PCA

```
In [ ]: model = []
# Model names
accuracy = []
# Accuracy of the respective model
```

```
In [ ]: from sklearn.decomposition import PCA
import sklearn.metrics as metrics
from sklearn.metrics import classification_report
pca = PCA(n_components=30)
htrain_pca = pca.fit_transform(htrain)

htest_pca=pca.transform(htest)

clf = svm.SVC(kernel='rbf') # rbf Kernel
#Train the model using the training sets
clf.fit(htrain_pca,train_y)
y_pred = clf.predict(htest_pca)
#Predict the response for test dataset
x = metrics.accuracy_score(test_y, y_pred)
model.append('SVM')
accuracy.append(x*100)
print(classification_report(test_y, y_pred))
print("SVM's Accuracy is: ", x*100)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	331
1	1.00	0.95	0.97	432
2	1.00	0.94	0.97	310
3	0.98	0.93	0.95	245
4	0.99	1.00	0.99	498
5	0.96	1.00	0.98	247
6	0.98	0.92	0.95	348
7	0.98	0.98	0.98	436
8	0.89	0.99	0.94	288
10	0.99	0.87	0.93	331
11	1.00	1.00	1.00	209
12	0.89	0.84	0.87	394
13	0.79	0.88	0.84	291
14	0.92	0.91	0.92	246
15	1.00	1.00	1.00	347
16	0.89	1.00	0.94	164
17	0.55	0.99	0.70	144
18	0.93	0.88	0.91	246
19	0.91	0.91	0.91	248
20	0.98	0.72	0.83	266
21	0.92	0.81	0.86	346
22	0.69	0.94	0.79	206
23	0.88	0.91	0.89	267
24	1.00	0.88	0.94	332
accuracy			0.93	7172
macro avg	0.92	0.93	0.92	7172
weighted avg	0.94	0.93	0.93	7172

SVM's Accuracy is: 92.63803680981594

```
In [ ]: y_pred = clf.predict(htrain_pca)
print("Training Set Metrics : ")
print("Accuracy:",metrics.accuracy_score(train_y, y_pred))
print("F1 Score:",metrics.f1_score(train_y, y_pred, average='weighted'))
print("Precision: ",metrics.precision_score(train_y,y_pred, average='weighted'))
)
print("Recall: ",metrics.recall_score(train_y,y_pred,average = 'weighted'))
```

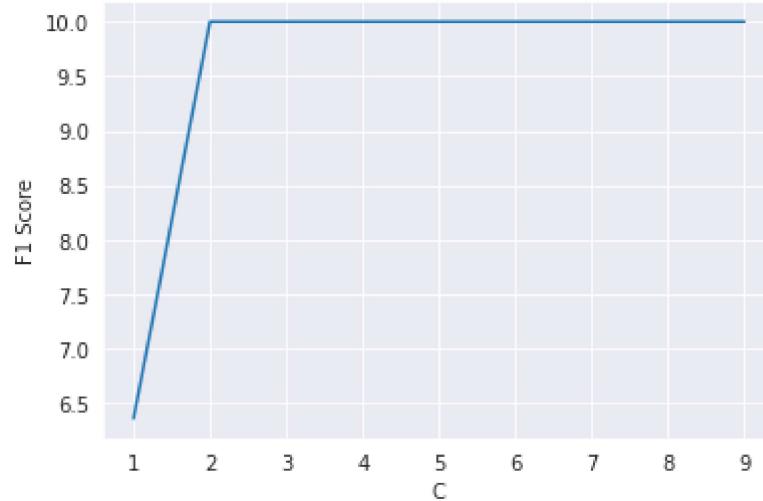
```
Training Set Metrics :
Accuracy: 0.9999635767619741
F1 Score: 0.9999635761187055
Precision: 0.9999636101471603
Recall: 0.9999635767619741
```

```
In [ ]: x=[]
y=[]
y2=[]
for i in range(1,10):
    clf = svm.SVC(kernel='rbf',C=i)
    clf.fit(htrain_pca,train_y)
    y_pred = clf.predict(htest_pca)
    x.append(i)
    y.append(metrics.f1_score(test_y, y_pred, average='weighted'))
    y2_pred = clf.predict(htrain_pca)
    y2.append(metrics.f1_score(train_y, y2_pred, average='weighted'))
```

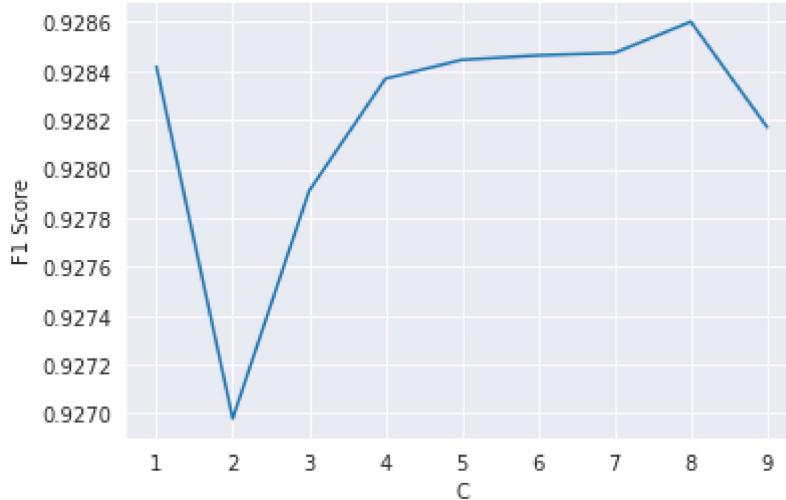
```
In [ ]: plt.plot(x,y2)
plt.title("F1 Score as a function of C(Regularization Parameter)) for training
          data")
plt.xlabel("C")
plt.ylabel("F1 Score")
plt.show()

plt.plot(x,y)
plt.title("F1 Score as a function of C(Regularization Parameter)) for test dat
          a")
plt.xlabel("C")
plt.ylabel("F1 Score")
plt.show()
```

F1 Score as a function of C(Regularization Parameter) for training data



F1 Score as a function of C(Regularization Parameter) for test data



Testing The Model

- For testing we are going to use the test data only

```
In [ ]: print('True:', test_y[25])
print('Pred:', y_pred[25])
```

True: 15
Pred: 15

Perceptron:

```
In [ ]: tran_y=pd.DataFrame(train_y)
```

```
In [ ]: tran_y.tail()
```

Out[]:

	0
27450	13
27451	23
27452	18
27453	17
27454	23

```
In [ ]: tes_y=pd.DataFrame(test_y)
```

```
In [ ]: tes_x=pd.DataFrame(test_x)
```

```
In [ ]: tran_x=pd.DataFrame(train_x)
```

```
In [ ]: from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score
clf = MLPClassifier()#solver='lbfgs', alpha=1e-5, hidden_layer_sizes=(5, 2), random_state=1
clf.fit(tran_x, tran_y)
y_pred=clf.predict(tes_x)
x = metrics.accuracy_score(tes_y, y_pred)
accuracy.append(x*100)
model.append('Perceptron')
print(classification_report(tes_y, y_pred))
print("Perceptron's Accuracy is: ", x*100)
```

	precision	recall	f1-score	support
0	0.84	1.00	0.92	331
1	0.99	0.91	0.94	432
2	0.90	0.91	0.91	310
3	0.84	0.82	0.83	245
4	0.94	0.91	0.93	498
5	0.64	0.91	0.75	247
6	0.85	0.81	0.83	348
7	0.87	0.72	0.79	436
8	0.78	0.75	0.76	288
10	0.70	0.49	0.58	331
11	0.73	0.77	0.75	209
12	0.76	0.79	0.78	394
13	0.84	0.57	0.68	291
14	0.95	0.63	0.76	246
15	1.00	0.97	0.98	347
16	0.58	0.78	0.67	164
17	0.24	0.43	0.31	144
18	0.43	0.48	0.45	246
19	0.47	0.53	0.50	248
20	0.38	0.49	0.43	266
21	0.75	0.60	0.67	346
22	0.64	0.70	0.67	206
23	0.58	0.68	0.62	267
24	0.77	0.64	0.70	332
accuracy			0.74	7172
macro avg	0.73	0.72	0.72	7172
weighted avg	0.76	0.74	0.74	7172

Perceptron's Accuracy is: 73.96820970440602

Logistic Regression:

Multi class Logistic Regression Using OVR

- Since we are going to use One Vs Rest algorithm, set > multi_class='ovr'
- Note: since we are using One Vs Rest algorithm we must use 'liblinear' solver with it.

```
In [ ]: from sklearn.linear_model import LogisticRegression
# train a logistic regression model on the training set
from sklearn import metrics
from sklearn import svm
from sklearn.metrics import classification_report

LogReg = LogisticRegression(multi_class='ovr', solver='liblinear')
# instantiate model
LogReg.fit(htrain,train_y)
# fit model
predicted_values = LogReg.predict(htest)
# make class predictions for the testing set
x = metrics.accuracy_score(test_y, predicted_values)
model.append('Logistic Regression')
accuracy.append(x*100)
print(classification_report(test_y, predicted_values))
print("Logistic Regression's Accuracy is: ", x*100)
```

	precision	recall	f1-score	support
0	0.99	1.00	0.99	331
1	0.97	0.90	0.93	432
2	0.95	0.93	0.94	310
3	0.81	0.99	0.89	245
4	0.95	1.00	0.97	498
5	0.95	0.97	0.96	247
6	0.82	0.85	0.83	348
7	0.96	0.94	0.95	436
8	0.90	0.87	0.88	288
10	0.91	0.74	0.81	331
11	0.92	1.00	0.96	209
12	0.79	0.62	0.70	394
13	0.75	0.65	0.70	291
14	0.78	0.79	0.78	246
15	0.97	1.00	0.99	347
16	0.76	0.87	0.81	164
17	0.37	0.64	0.47	144
18	0.70	0.80	0.74	246
19	0.71	0.85	0.77	248
20	0.81	0.61	0.70	266
21	0.78	0.56	0.65	346
22	0.58	0.91	0.71	206
23	0.71	0.72	0.71	267
24	0.94	0.75	0.83	332
accuracy			0.84	7172
macro avg	0.82	0.83	0.82	7172
weighted avg	0.85	0.84	0.84	7172

Logistic Regression's Accuracy is: 83.63078639152258

Random Forest:

```
In [ ]: from sklearn.ensemble import RandomForestClassifier

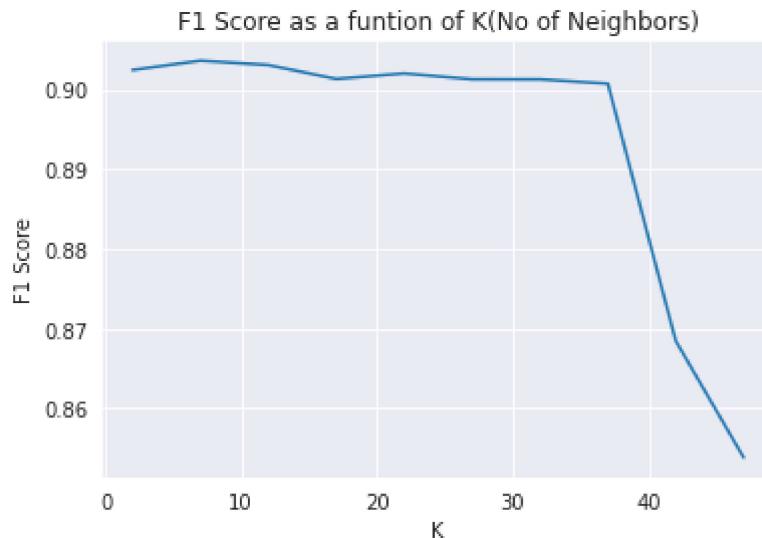
RF = RandomForestClassifier()
RF.fit(train_x,train_y)
predicted_values = RF.predict(test_x)
x = metrics.accuracy_score(test_y, predicted_values)
accuracy.append(x*100)
model.append('Random Forest')
print(classification_report(test_y,predicted_values))
print("RF's Accuracy is: ", x*100)
```

	precision	recall	f1-score	support
0	0.91	1.00	0.96	331
1	0.96	0.94	0.95	432
2	0.94	0.98	0.96	310
3	0.81	0.97	0.88	245
4	0.89	0.97	0.93	498
5	0.91	0.93	0.92	247
6	0.92	0.86	0.89	348
7	0.99	0.92	0.95	436
8	0.81	0.70	0.75	288
10	0.72	0.62	0.67	331
11	0.81	1.00	0.90	209
12	0.85	0.67	0.75	394
13	0.73	0.58	0.65	291
14	0.97	0.87	0.92	246
15	0.95	1.00	0.97	347
16	0.91	1.00	0.95	164
17	0.29	0.56	0.38	144
18	0.60	0.81	0.69	246
19	0.61	0.82	0.70	248
20	0.65	0.63	0.64	266
21	0.82	0.58	0.68	346
22	0.55	0.67	0.60	206
23	0.86	0.69	0.77	267
24	0.86	0.61	0.71	332
accuracy			0.81	7172
macro avg	0.80	0.81	0.80	7172
weighted avg	0.83	0.81	0.82	7172

RF's Accuracy is: 81.39988845510318

KNN

```
In [ ]: from sklearn.neighbors import KNeighborsClassifier
x=[]
y=[]
for i in range(2,50,5):
    neigh = KNeighborsClassifier(n_neighbors=i)
    neigh.fit(htrain_pca,train_y)
    y_pred = neigh.predict(htest_pca)
    x.append(i)
    y.append(metrics.f1_score(test_y, y_pred, average='weighted'))
plt.plot(x,y)
plt.title("F1 Score as a function of K(No of Neighbors)")
plt.xlabel("K")
plt.ylabel("F1 Score")
plt.show()
```



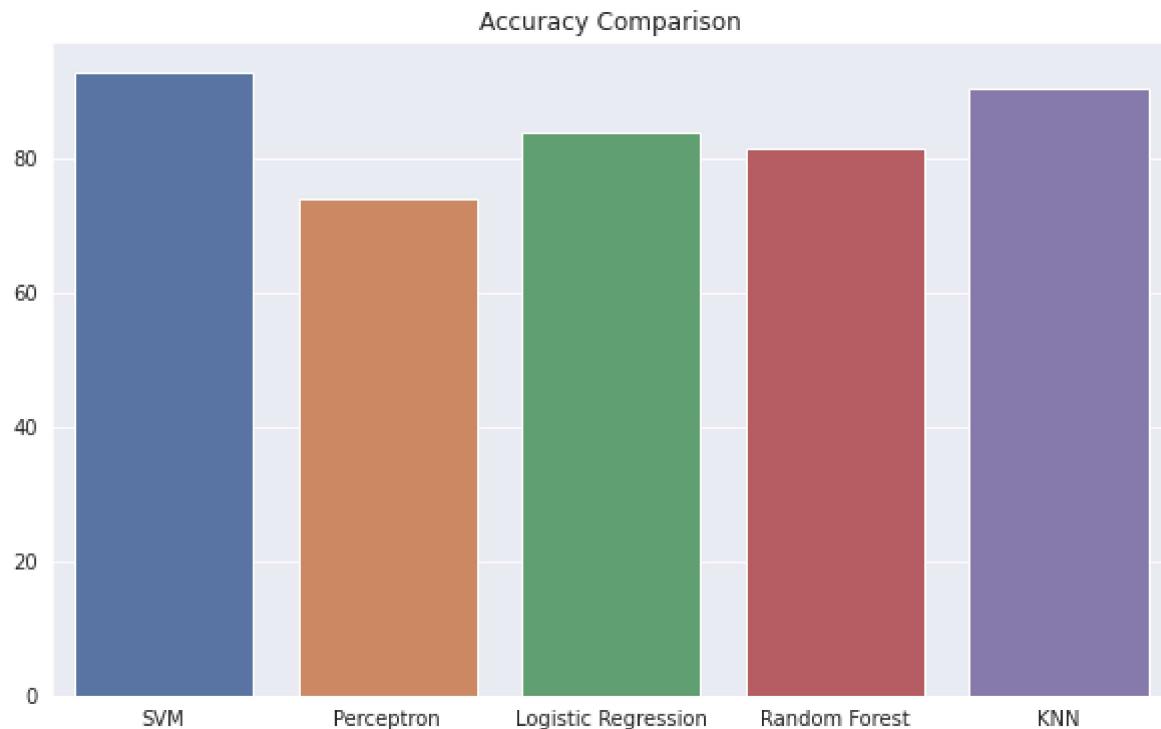
```
In [ ]: m=max(y)
index = y.index(m)
print("Max value of F1 Score occurs at K = ",x[index])
print("The Evaluation metrics are : ")
neigh = KNeighborsClassifier(n_neighbors=x[index])
neigh.fit(htrain_pca,train_y)
y_pred = neigh.predict(htest_pca)
x = metrics.accuracy_score(test_y, y_pred)
accuracy.append(x*100)
model.append('KNN')
print("Accuracy:",metrics.accuracy_score(test_y, y_pred))
print("F1 Score:",metrics.f1_score(test_y, y_pred, average='weighted'))
print("Precision: ",metrics.precision_score(test_y,y_pred, average='weighted'))
print("Recall: ",metrics.recall_score(test_y,y_pred,average = 'weighted'))
```

Max value of F1 Score occurs at K = 7
The Evaluation metrics are :
Accuracy: 0.9033742331288344
F1 Score: 0.9036218462600369
Precision: 0.9104427818377326
Recall: 0.9033742331288344

Model Comparison

```
In [ ]: plt.figure(figsize=(10,6))
plt.title('Accuracy Comparison')
sns.barplot(x=model, y=accuracy, palette='deep')
```

```
Out[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fbb03a11090>
```



```
In [ ]: for x in range(len(accuracy)):
    print(accuracy[x])
```

```
92.63803680981594
73.96820970440602
83.63078639152258
81.39988845510318
90.33742331288343
```

```
In [ ]: def Average(accuracy):
    return sum(accuracy) / len(accuracy)
average = Average(accuracy)
print("Average accuracy of classifier models is =", round(average, 2))
```

```
Average accuracy of classifier models is = 84.39
```

To visualize the inner working of the model we use Confusion Matrix

```
In [ ]: y_predicted = neigh.predict(htest_pca)
```

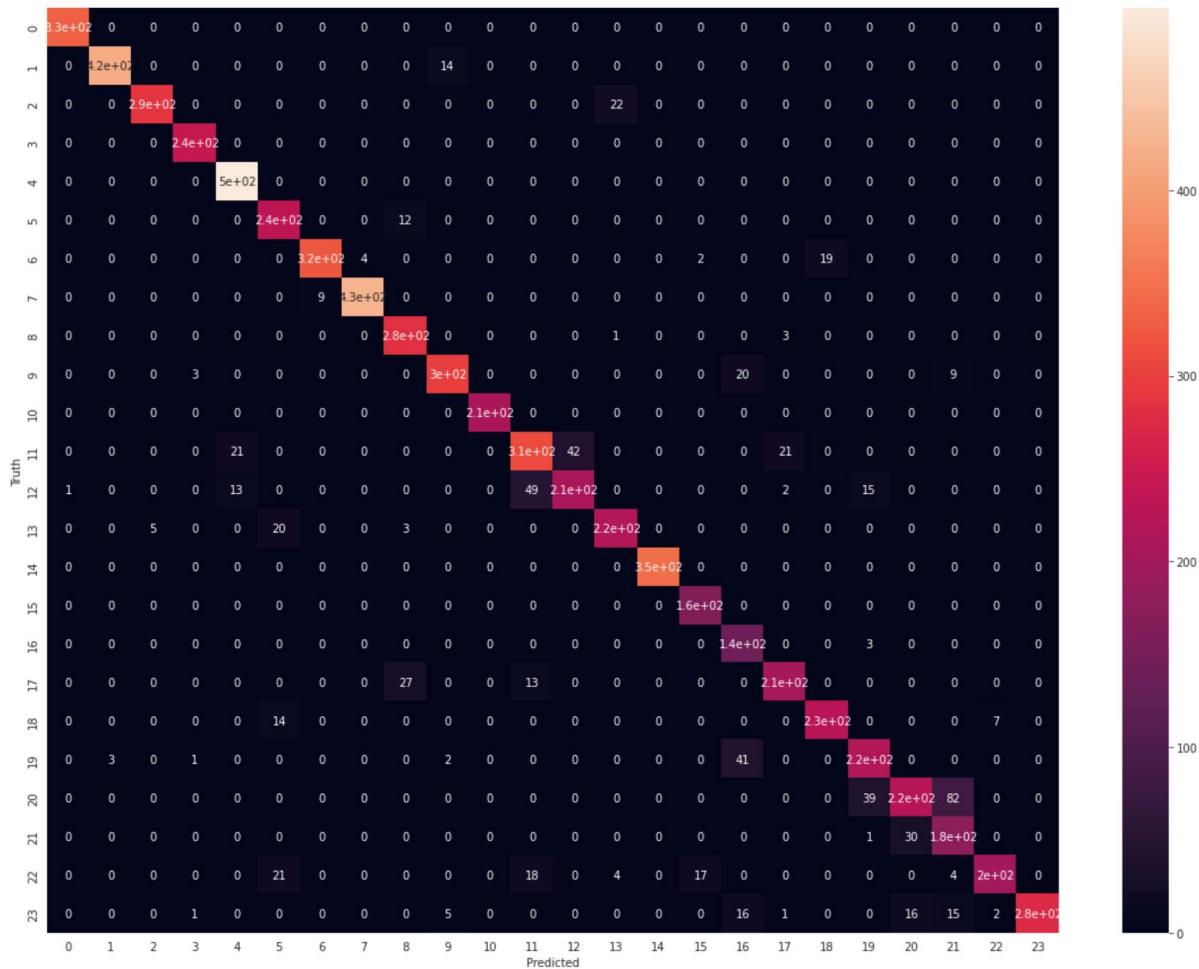
```
In [ ]: print('True:', test_y[0:25])
# print the first 25 true and predicted responses
print('Pred:', y_predicted[0:25])
```

```
True: [ 6  5 10  0  3 21 10 14  3  7  8  8 21 12  7  4 22  0  7  7  2  0 21
4
10]
Pred: [ 6  5 10  0  3 20 10 14  3  7  8  8 22 12  7  4 22  0  7  7  2  0 21
4
10]
```

```
In [ ]: from sklearn.metrics import confusion_matrix  
# import confusion matrix class  
cm = confusion_matrix(test_y,y_predicted)  
# create classifier of confusion_matrix  
cm  
# it is two dimentional arraym
```

```
In [ ]: plt.figure(figsize = (20,15))
sns.heatmap(cm, annot=True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Out[]: Text(159.0, 0.5, 'Truth')



Observation:

- In above matrix, X axis contains predicted values , and Y axis contains actual values .
- Each value of matrix show the number of times predicted value matched with actual value.

Cross Validation:

```
In [ ]: X_train = train_df[train_df.columns[1:]]
y_train = train_df[train_df.columns[0]]

X_test = test_df[test_df.columns[1:]]
y_test = test_df[test_df.columns[0]]
```

```
In [ ]: X_train=X_train/255  
X_test=X_test/255
```

```
In [ ]: from sklearn import svm  
clf = svm.SVC(kernel='rbf', C=2).fit(X_train, y_train)  
clf.score(X_test, y_test)
```

```
In [ ]: from sklearn.model_selection import cross_val_score  
scores = cross_val_score(clf, X_train, y_train, cv=5)  
scores
```

```
In [ ]: print("%0.2f accuracy with a standard deviation of %0.2f" % (scores.mean(),  
scores.std()))
```

Got a mean of 0.98 with std=0.2

CONCLUSIONS AND FUTURE WORK

Hand gesture recognition is a difficult problem and the current work is only half way towards achieving the results needed in the field.

- This notebook presented a comparative study of five algorithms applied to two different datasets for static gesture recognition and classification, for human computer interaction.
- Future work will be concerned with the study of different hand feature selection applied to hand gesture recognition, noise reduction in the depth images acquired.

References:

1. https://en.wikipedia.org/wiki/Support-vector_machine (https://en.wikipedia.org/wiki/Support-vector_machine)
2. <https://www.datacamp.com/community/tutorials/svm-classification-scikit-learn-python> (<https://www.datacamp.com/community/tutorials/svm-classification-scikit-learn-python>)
3. <https://link.springer.com/article/10.1007/s10586-021-03282-8> (<https://link.springer.com/article/10.1007/s10586-021-03282-8>).